

Verifiable Confidential Token Contract

Roman Walch¹

¹ TACEO

papers@taceo.io

Abstract. In this paper we describe how to build a confidential token contract with MPC and collaborative SNARKs. This smart contract hides the balances of users, as well as transaction amounts when sending tokens to other participants by only showing commitments publicly. The actual balances corresponding to these commitments are stored – as secret shares – on an MPC network, which performs the necessary computations privately in MPC. The addition of the collaborative SNARK adds verifiability to the system, proving that only valid state transitions are performed and enough funds are present for each transaction. Our system currently achieves a throughput of 200 transactions per second (including proof generation) on reasonable MPC hardware.

Keywords: MPC · ZKP · CoSNARKs · Confidential Token

Contents

1	Introduction	2
1.1	Related Work	2
1.2	Notation	3
2	High Level Overview	4
3	The MPC Computations	4
4	The Zero Knowledge Proofs	6
4.1	Deposit Proof π_{deposit}	6
4.2	Withdraw Proof π_{withdraw}	7
4.3	Transfer Proof π_{transfer}	8
4.4	Batching Proofs	9
5	The Full Protocol	10
5.1	Read Balance	10
5.2	Deposit	11
5.3	Withdraw	11
5.4	Transfer	12
6	Benchmarks	14
6.1	Single Benchmarks	14

6.2	Batched Benchmarks	15
6.3	Solidity Gas Costs	16
7	Issues and Future Work	17
7.1	Implications of Batching Transactions	17
7.2	Concurrent Transactions and Sharding	18
7.3	Hiding Involved Parties	18
	References	18

1 Introduction

Privacy preserving cryptographic protocols and primitives such as secure multiparty computation (MPC) [Yao82], fully homomorphic encryption (FHE) [Gen09] and zero-knowledge proofs (ZKP) have been heavily improved in the recent years and become more feasible to apply to real world use cases every day. In this document we investigate the application of these protocols, more concretely MPC, ZKPs and their intersection, dubbed collaborative SNARKs (CoSNARKs) [OB22], to the use case of private on-chain payments. Concretely, we investigate the case of confidential and verifiable on-chain token balances. There, the goal is to add privacy to an existing token (which can for example be ETH, or a stable coin such as USDC), by adding a smart contract which stores balances of users privately, e.g., by holding cryptographic commitments or (homomorphic) ciphertexts of the balances. Then, when users transfer tokens amongst each other, the commitments or ciphertexts are updated accordingly, while zero knowledge proofs prove that 1) the sender has enough balance and 2) the commitments were updated correctly. In this way, the current balance of the involved users, as well as the transferred amount, stay private while everyone can verify the computations were done correctly.

The users should also be able to deposit fresh tokens to its private on-chain balance, and withdraw them (if enough funds are present). To this end, when the user deposits some tokens into the smart contract, the smart contract keeps it alongside all other deposited tokens in a pool. Furthermore, when the user withdraws funds, the smart contract pays out tokens from this pool. Thus, the amount of tokens in this pool should always match the total balance of all users in the system.

1.1 Related Work

Private onchain payments have seen many different approaches in the literature, which we review here shortly.

First, approaches based purely on ZKPs have been studied and deployed in practice (e.g., Zexe [Bow+20], Zcash¹). The main approach here is that whenever parties want to transact privately, they create a so-called note which can only be decrypted and spent by a dedicated receiver. Due to employing commitments and ZK proofs, on-chain data does not leak any information on the balances and transaction amounts stays hidden. When a party spends a

¹<https://z.cash/>

note, it releases a nullifier, which prevents the note from being spent a second time. The main drawback of these systems, however, is that users do not get to know that they received some notes without constantly scanning and trying to decrypt all notes to determine which belong to themselves.

The possibility to store and modify encrypted data (as given by FHE and MPC) allows a third party to store the actual balances in encrypted form without leaking the actual data. On request, the third party opens the balance to be viewed by a user, mitigating this note discovery problem. While we propose an MPC based solution in this document, solutions based on homomorphic encryption were already proposed, which we discuss in the following paragraph.

The main idea in HE-based solutions is that one server holds a ciphertexts of each users balances, while the corresponding decryption key is secret shared amongst an MPC network. To read balances, a client requests a decryption from the MPC network, which engages in a MPC protocol (e.g., as proposed in [Zys+25]) to output the balance in plaintext to the user. When requesting a transfer, the client sends ciphertexts to the server holding the balance, which can modify the current balances accordingly. Assuming the ciphertexts are well-formed, the whole protocol is public-transcript. In other words, clients can redo computations and challenge the server if it performed faulty computation, leading to an *optimistic* verifiability property. To this end, clients should provide a zero knowledge proof, proving that the ciphertexts they send are well-formed.

The main drawback of these HE solutions is performance. First, threshold decryption is rather slow ([Zys+25] gives a throughput of 705 decryptions per second for 8-bit ciphertexts in a 2 party setup, while ignoring precomputations of Beaver triples [Bea91]), and lack verifiability in most deployments.² Second, to achieve high throughput, expensive server farms are required. E.g, TFHE [Chi+20], as implemented by Zama, achieves a throughput of 635 64-bit integer additions per second³ on a server farm with 8x H100 GPUs, while a 96-core AWS hpc7a.96xlarge instance only achieves a throughput of 72.5 additions per second⁴. Finally, client side proofs are also expensive, requiring the client to compute for 1.5s when having access to a 16 core machine.⁵

Finally, we briefly want to mention that approaches based on trusted execution environments (TEEs) have also been proposed in the literature. In a sense, they work very similar to the HE solution outlined above, with the exception that the balances are stored in a TEE instead of being encrypted. This approach trades improved performance against significant lower privacy guarantees, since TEEs are known to be a prime target for side channel attacks.⁶

1.2 Notation

In the rest of the manuscript, we denote a secret sharing of a value $x \in \mathbb{F}_p$ with a bracket notation $[x]$.

²Extending the MPC computation with a CoSNARK could add verifiability at the cost of less throughput.

³<https://docs.zama.ai/tfhe-rs/get-started/benchmarks/gpu/gpu-integer-operations>

⁴<https://docs.zama.ai/tfhe-rs/get-started/benchmarks/cpu/cpu-integer-operations>

⁵<https://docs.zama.ai/tfhe-rs/get-started/benchmarks/zk-proof-benchmarks>

⁶See, e.g., <https://sgx.fail/> for a summary of attacks.

2 High Level Overview

In this section we describe our proposed solution on a high level, for more details on the involved subprotocols we refer to the next sections.

In our solution, we essentially have three different entities involved:

- A user who deposits and withdraws tokens and wants to initiate private token transfers to others.
- A smart contract keeping private balances and a pool of deposited tokens.
- An MPC-network holding secret-shares of the actual balances.

Keeping not only transaction amounts and balances private, but also which wallets interact with each other will lead to significantly more MPC workload (and thus less performance) and might give raise to concerns around compliance topics. Hence we do not hide sender/receiver addresses and the smart contract holds a hashmap (e.g, mapping in solidity), mapping user addresses directly to commitments. Furthermore, the MPC network holds another hashmap, which maps user addresses to secret shares of the balance and the randomness used to compute the commitment which is stored on chain. Whenever the user initiates an action (i.e., deposit, withdraw, or transfer), it has to communicate this to the smart contract. The smart contract will then register the action in a queue which is continuously monitored by the MPC network. The MPC network then computes the action in MPC which leads to updated commitments. To prove it performed the updates correctly, it also generates a zero knowledge proof in MPC (which is called collaborative SNARK, or CoSNARK) proving correctness of the update. The MPC network posts the new commitments alongside the proof on chain, where the smart contract verifies the proof, updates the commitments in its hashmap if the proof was correct, removes the action from the queue, and, in case of withdraw, sends tokens to the recipient.

The advantages of this systems are as follows. First, the client side computations are minimal. To initiate transactions, it only has to post a transaction on chain and, in case of a transfer, send secret shares to the MPC network. No expensive client side ZK proofs need to be computed. Second, since the MPC network creates ZK proofs of all its actions, integrity of funds is always given and the MPC network can not misuse funds.

3 The MPC Computations

In this section we have a closer look at the computations happening at the MPC network. We list all the involved algorithms in Algorithm 1. The first two algorithms in Algorithm 1 (`Map.insert` and `Map.get`) represent the access of the actual data structure, which store the secret-shared balance and randomness. Note that these are normal hashmaps. Then, each action (deposit, withdraw and transfer) essentially just create new commitment and stores them in the hashmap. Finally each action concludes by computing a zero knowledge proof in MPC to prove integrity to the smart contract.

Algorithm 1. The computations on the MPC network

```

1:  ▷ Data structure
2:  function MAP.INSERT( $k$ ,  $[b]$ ,  $[r]$ )
3:    ▷ Store balance  $[b]$  and randomness  $[r]$  in the hashmap at key  $k$ .
4:  end
5:
6:  function MAP.GET( $k$ )
7:    if hashmap has entry( $[b]$ ,  $[r]$ ) for key  $k$  then
8:      return ( $[b]$ ,  $[r]$ )
9:    else
10:      return (0, 0)
11:    end
12:  end
13:
14: function DEPOSIT( $k$ ,  $b$ )
15:    $[r_{\text{old}}], [b_{\text{old}}] \leftarrow \text{Map.get}(k)$ 
16:    $[r_{\text{new}}] \leftarrow \text{MPC.rand}()$                                 ▷ Sample new blinding for commitment
17:    $[b_{\text{new}}] \leftarrow [b_{\text{old}}] + b$ 
18:    $\text{Map.insert}(k, [b_{\text{new}}], [r_{\text{new}}])$ 
19:    $C_{\text{old}} \leftarrow \text{commit}([b_{\text{old}}], [r_{\text{old}}])$ 
20:    $C_{\text{new}} \leftarrow \text{commit}([b_{\text{new}}], [r_{\text{new}}])$ 
21:    $\pi_{\text{deposit}} \leftarrow \text{proof}(C_{\text{old}}, C_{\text{new}}, b)$                 ▷ Collaborative SNARK, see Listing 1
22:   return  $\pi_{\text{deposit}}, C_{\text{new}}$ 
23: end
24:
25: function WITHDRAW( $k$ ,  $b$ )
26:    $[r_{\text{old}}], [b_{\text{old}}] \leftarrow \text{Map.get}(k)$ 
27:    $[r_{\text{new}}] \leftarrow \text{MPC.rand}()$                                 ▷ Sample new blinding for commitment
28:    $[b_{\text{new}}] \leftarrow [b_{\text{old}}] - b$ 
29:    $\text{Map.insert}(k, [b_{\text{new}}], [r_{\text{new}}])$ 
30:    $C_{\text{old}} \leftarrow \text{commit}([b_{\text{old}}], [r_{\text{old}}])$ 
31:    $C_{\text{new}} \leftarrow \text{commit}([b_{\text{new}}], [r_{\text{new}}])$ 
32:    $\pi_{\text{withdraw}} \leftarrow \text{proof}(C_{\text{old}}, C_{\text{new}}, b)$                 ▷ Collaborative SNARK, see Listing 2
33:   return  $\pi_{\text{withdraw}}, C_{\text{new}}$ 
34: end
35:
36: function TRANSFER( $k_{\text{sender}}$ ,  $k_{\text{receiver}}$ ,  $[b]$ ,  $[r_b]$ )
37:    $[r_{s,\text{old}}], [b_{s,\text{old}}] \leftarrow \text{Map.get}(k_{\text{sender}})$ 
38:    $[r_{s,\text{new}}] \leftarrow \text{MPC.rand}()$                                 ▷ Sample new blinding for commitment
39:    $[b_{s,\text{new}}] \leftarrow [b_{s,\text{old}}] - [b]$ 
40:    $\text{Map.insert}(k_{\text{sender}}, [b_{s,\text{new}}], [r_{s,\text{new}}])$ 
41:    $C_{s,\text{old}} \leftarrow \text{commit}([b_{s,\text{old}}], [r_{s,\text{old}}])$ 
42:    $C_{s,\text{new}} \leftarrow \text{commit}([b_{s,\text{new}}], [r_{s,\text{new}}])$ 
43:
44:    $[r_{r,\text{old}}], [b_{r,\text{old}}] \leftarrow \text{Map.get}(k_{\text{receiver}})$ 
45:    $[r_{r,\text{new}}] \leftarrow \text{MPC.rand}()$                                 ▷ Sample new blinding for commitment

```

```

46:    $[b_{r,new}] \leftarrow [b_{r,old}] \ddot{u}[b]$ 
47:   Map.insert( $k_{receiver}$ ,  $[b_{r,new}]$ ,  $[r_{r,new}]$ )
48:    $C_{r,old} \leftarrow \text{commit}([b_{r,old}], [r_{r,old}])$ 
49:    $C_{r,new} \leftarrow \text{commit}([b_{r,new}], [r_{r,new}])$ 
50:
51:    $C_b \leftarrow \text{commit}([b], [r_b])$ 
52:    $\pi_{\text{transfer}} \leftarrow \text{proof}(C_{s,old}, C_{s,new}, C_{r,old}, C_{r,new}, C_b)$   $\triangleright$  Collaborative SNARK, see Listing 3
53:   return  $\pi_{\text{transfer}}, C_{s,new}, C_{r,new}$ 
54: end

```

4 The Zero Knowledge Proofs

In this section we give details to the zero knowledge proof created by the MPC network (see Algorithm 1).

4.1 Deposit Proof π_{deposit}

The deposit proof (Listing 1) essentially proves that the commitment of the targeted user has been updated correctly. It thus can be verified using the old commitment (stored in the smart contract), the new commitment (is put into the smart contract by the MPC network) and by knowing the deposit amount (which is also known by the smart contract at this point). Furthermore, we limit the deposit amount to a reasonable number, which is also part of the ZK proof. The reason for this limit is that withdraw and transfer proofs need to show that the new balance of the sender is still positive. Having a limit on the transfer amount simplifies this range proof.

In our current prototype we use a limit of 80 bits for transaction amounts. In case of ethereum where balances are denoted in Wei, this allows to transact $\frac{2^{80}}{10^{18}} \approx 1.2$ million ETH at once, which should be enough for most use cases.

Proof $\pi_{\text{deposit}} \leftarrow \text{prove}(C_{\text{old}}, C_{\text{new}}, b)$:

- Private witnesses:
 - Old balance b_{old}
 - Old blinding r_{old}
 - New blinding r_{new}
- Public witnesses:
 - Amount deposit b
 - Old commitment C_{old}
 - New commitment C_{new}
- Statements:
 - $b_{\text{new}} \leftarrow b_{\text{old}} + b$
 - C_{old} equals $\text{commit}(b_{\text{old}}, r_{\text{old}})$
 - C_{new} equals $\text{commit}(b_{\text{new}}, r_{\text{new}})$
 - b is in range (e.g., 80 bits)

Listing 1. The statements proven in π_{deposit} .

4.2 Withdraw Proof π_{withdraw}

The withdraw proof (Listing 2) is conceptually very similar to the deposit proof (Listing 1), with the difference that the balance is subtracted instead of added, and that the resulting balance proven to be ≥ 0 . The latter statement is proven by a range check.

In our current prototype we somewhat arbitrarily prove that the resulting statement has at most 100 bits.

Proof $\pi_{\text{transfer}} \leftarrow \text{prove}(C_{\text{old}}, C_{\text{new}}, b)$:

- Private witnesses:
 - Old balance b_{old}
 - Old blinding r_{old}
 - New blinding r_{new}
- Public witnesses:
 - Amount withdrawn b
 - Old commitment C_{old}
 - New commitment C_{new}
- Statements:
 - $b_{\text{new}} \leftarrow b_{\text{old}} - b$
 - C_{old} equals $\text{commit}(b_{\text{old}}, r_{\text{old}})$
 - C_{new} equals $\text{commit}(b_{\text{new}}, r_{\text{new}})$
 - b is in range (e.g., 80 bits)
 - b_{new} is zero or positive (e.g., by checking it is in range of, e.g., 100 bits)

Listing 2. The statements proven in π_{withdraw} .

4.3 Transfer Proof π_{transfer}

Finally, the transfer proof (Listing 3) is essentially a combination of the deposit proof (Listing 1) and withdraw proof (Listing 2). In other words, the proof shows that the sender commitments were updated as in a withdraw (including proof of enough balances), whereas the receiver commitments were updated according to a deposit proof.

Proof $\pi_{\text{withdraw}} \leftarrow \text{prove}(C_{s,\text{old}}, C_{s,\text{new}}, C_{r,\text{old}}, C_{r,\text{new}}, C_b)$:

- Private witnesses:
 - Old sender balance $b_{s,\text{old}}$
 - Old sender blinding $r_{s,\text{old}}$
 - New sender blinding $r_{s,\text{new}}$
 - Old receiver balance $b_{r,\text{old}}$
 - Old receiver blinding $r_{r,\text{old}}$
 - New receiver blinding $r_{r,\text{new}}$
 - Amount transferred b
 - Amount blinding r_b
- Public witnesses:
 - Old sender commitment $C_{s,\text{old}}$
 - New sender commitment $C_{s,\text{new}}$
 - Old receiver commitment $C_{r,\text{old}}$
 - New receiver commitment $C_{r,\text{new}}$
 - Amount commitment C_b
- Statements:
 - $b_{s,\text{new}} \leftarrow b_{s,\text{old}} - b$
 - $b_{r,\text{new}} \leftarrow b_{r,\text{old}} + b$
 - $C_{s,\text{old}}$ equals $\text{commit}(b_{s,\text{old}}, r_{s,\text{old}})$
 - $C_{s,\text{new}}$ equals $\text{commit}(b_{s,\text{new}}, r_{s,\text{new}})$
 - $C_{r,\text{old}}$ equals $\text{commit}(b_{r,\text{old}}, r_{r,\text{old}})$
 - $C_{r,\text{new}}$ equals $\text{commit}(b_{r,\text{new}}, r_{r,\text{new}})$
 - C_b equals $\text{commit}(b, r_b)$
 - b is in range (e.g., 80 bits)
 - $b_{s,\text{new}}$ is positive (e.g., by checking it is in range of, e.g., 100 bits)

Listing 3. The statements proven in π_{transfer} .

4.4 Batching Proofs

Since verifying a ZK prove on-chain is costly in terms of gas fees (mostly due to evaluating a pairing), we investigate batching multiple transfers into one ZK proof. This would reduce the gas fee for verifying multiple transfers significantly. Furthermore, `deposit` and `withdraw` can be proven by π_{transfer} as well, allowing to batch deposits, withdrawals and individual transfers into just one ZK proof which verifies multiple transfers.

To verify a `deposit` via a transfer proof π_{transfer} , we do the following:

- Use the old balance, old blinding and new blinding of the `deposit` as the old receiver balance, old receiver blinding and new receiver blinding in a `transfer` respectively.

- Set the amount blinding to 0.
- Set the old sender balance to the amount b , the old sender blinding to 0, and the new sender blinding to 0 as well.

To verify a `withdraw` via a transfer proof π_{transfer} , we do the following:

- Use the old balance, old blinding and new blinding of the `withdraw` as the old sender balance, old sender blinding and new sender blinding in a `transfer` respectively.
- Set the amount blinding to 0.
- Set the old receiver balance to 0, the old receiver blinding to 0, and the new receiver blinding to the amount b .

5 The Full Protocol

In this section section, we describe the full protocols for reading balances, depositing/withdrawing funds and privately transferring tokens.

5.1 Read Balance

In our proposed system, a user can read balances by asking the MPC network to send over the balance and the randomness used for the commitment in plain. Notice that no zero knowledge proof needs to be involved, since the user can just read the commitment corresponding to itself from the smart contract and verify whether it matches what it received from the MPC network by recomputing the commitment. Consequently, no smart contract transaction is involved as well, only a read of the commitment if a user wants to verify the response of the MPC network. Finally, to keep balances private, the user needs to send a signature to the MPC network, proving its identity. The full protocol is given in Listing 4.

- **The User:**
 1. The user signs a timestamp t to create the signature $\sigma \leftarrow \text{sign}(t, \text{sk})$ using its private key sk . The user sends (t, σ, pk) to the MPC network
- **The MPC Network:**
 2. On receiving (t, σ, pk) from a user, the MPC network verifies that the timestamp t is still valid, verifies the signature $\text{verify}(\sigma, t, \text{pk})$, and gets $([b], [r]) \leftarrow \text{Map.get}(\text{address}(\text{pk}))$. It then encrypts $[b]$ and $[r]$ for the user using pk and transmits the ciphertexts.
- **The User:**
 3. On receiving (b, r) from the MPC network, the user reads its commitment C from the smart contract and checks whether it matches $\text{commit}(b, r)$.

Listing 4. The protocol for reading balances.

5.2 Deposit

For a deposit, the user first has to transmit funds into the smart contract. Doing so, the smart contract registers a deposit action in a queue, which gets processed by the MPC network. While processing, the MPC network updates its shares corresponding to the user address, computes a commitment to the new balance, and a ZK proof of correctness. It then posts the proof on-chain, where the smart contract verifies the proof and stores the new commitment in its internal hashmap to show the user balance has been updated. The full protocol is given in Fig. 2.

Note that the deposit does not include any signatures, since the smart contract transaction implicitly authenticates the user already.

- **The User:**
 1. The user submits a transaction to the smart contract, sending b tokens to it.
- **The Smart Contract:**
 2. The smart contract registers a deposit action to its MPC queue, storing the users address k as receiver and the balance b as amount. If b is zero or too large it will revert.
- **The MPC Network:**
 3. The MPC network checks the queue on the smart contract and computes $(\pi_{\text{deposit}}, C_{\text{new}}) \leftarrow \text{Deposit}(k, b)$. It posts π_{deposit} and C_{new} into the smart contract.
- **The Smart Contract:**
 4. The smart contract receives π_{deposit} and C_{new} , verifies the proof π_{deposit} using the received C_{new} and the remaining public information which are computable on chain, sets C_{new} to be the new commitment corresponding to k and removes the action from the MPC queue.

Fig. 2. The protocol for depositing tokens.

5.3 Withdraw

For a withdraw, the user first has to post the intent into the smart contract. Doing so, the smart contract registers a withdraw action in a queue, which gets processed by the MPC network. While processing, the MPC network updates its shares corresponding to the user address, computes a commitment to the new balance, and a ZK proof proving correctness. It then posts the proof on chain, where the smart contract verifies the proof and stores the new commitment in its internal hashmap to show the user balance has been updated. If everything was done correctly, the smart contract sends funds to the user. The full protocol is given in Fig. 2.

Note that the withdraw does not include any signatures, since the smart contract transaction implicitly authenticates the user already. Furthermore, if the proof does not verify, the MPC network needs to roll back the transaction to keep secret shares of the old balance.

- **The User:**
 1. The user submits a transaction to the smart contract, indicating it wants to withdraw b tokens.
- **The Smart Contract:**
 2. The smart contract registers a withdraw action to its MPC queue, storing the users address k as sender and the balance b as amount. If b is zero or too large it will revert.
- **The MPC Network:**
 3. The MPC network checks the queue on the smart contract and computes $(\pi_{\text{withdraw}}, C_{\text{new}}) \leftarrow \text{Withdraw}(k, b)$. It posts π_{withdraw} and C_{new} into the smart contract.
- **The Smart Contract:**
 4. The smart contract receives π_{withdraw} and C_{new} , verifies the proof π_{withdraw} using the received C_{new} and the remaining public information which are computable on chain, sets C_{new} to be the new commitment corresponding to k and removes the action from the MPC queue. It then transmits b tokens to the user address k .

Fig. 3. The protocol for withdrawing tokens.

We note that it is possible to implement a fail-safe functionality into the smart contract that allows users to retrieve funds if the MPC network would go out of service. This could be realized by letting users withdraw their funds by providing the amount and the randomness which were used to compute the on-chain commitment. This fail-safe, however, should only be enabled on a trigger (e.g., if the MPC network does not make a keep-alive transaction for a longer time) to prevent the on-chain commitments going out of sync with the MPC network. Furthermore, this technique requires the user to regularly query the randomness and balance from the MPC network.

5.4 Transfer

When a user wants to transfer funds to someone else, it first needs to post the intent into the smart contract. This intent includes the receiver address, as well as a commitment to the amount which later gets linked to the zero knowledge proof posted by the MPC network. Then, the user needs to secret-share the transfer amount, as well as the randomness used to compute the commitments, to the MPC network. To authenticate the shares and prevent other users to post shares for transfers which are not theirs, the user has to sign the secret shares it sends to the MPC network. The MPC network, on reading the queue and receiving the secret shares,

updates its balances, computes commitments to the new balances, and computes a ZK proof proving correctness. It then posts the proof on chain, where the smart contract verifies the proof and stores the new commitments in its internal hashmap to show the user balances have been updated.

Note that if the proof does not verify, the MPC network needs to roll back the transaction to keep secret shares of the old balances.

The client can, instead of sending the secret shares to the MPC network, post encryptions of the shares on chain while registering the transfer. This trades the advantage of not requiring interaction between the MPC network and the client with a higher gas and calldata cost for the client when submitting the ciphertexts to the smart contract.

- **The User:**

1. The user samples a random value r_b and computes the commitment $C_b \leftarrow \text{commit}(b, r_b)$. It then submits a transaction to the smart contract, indicating it wants to transfer tokens to k_{receiver} , where the amount corresponds to C_b . It also sends the secret share b_i of b and the secret share r_i of r_b to the i -th MPC party, alongside its user address k_{sender} and a signature $\sigma_i \leftarrow \text{sign}(b_i \parallel r_i, pk)$, where pk corresponds to k_{sender} .

- **The Smart Contract:**

2. The smart contract registers a transfer action to its MPC queue, storing the users address k_{sender} as sender, the address k_{receiver} as receiver, and the commitment C_b .

- **The MPC Network:**

3. The MPC network receives $([b], [r_b])$ and the signatures σ_i for each party i . Each party i verifies the signature σ_i using its shares b_i and r_i , and matches the shares with the transfer request in the queue of the smart contract. If everything is ok, it computes $(\pi_{\text{transfer}}, C_{s,\text{new}}, C_{r,\text{new}}) \leftarrow \text{Transfer}(k_{\text{sender}}, k_{\text{receiver}}, [b], [r_b])$. It posts π_{transfer} and the commitments $C_{s,\text{new}}$ and $C_{r,\text{new}}$ into the smart contract.

- **The Smart Contract:**

4. The smart contract receives π_{transfer} and the commitments $C_{s,\text{new}}$ and $C_{r,\text{new}}$ and verifies the proof π_{transfer} using the received commitments and the remaining public information which are computable on chain. It then sets $C_{s,\text{new}}$ to be the new commitment corresponding to k_{sender} , $C_{r,\text{new}}$ to be the new commitment corresponding to k_{receiver} and removes the action from the MPC queue.

Fig. 4. The protocol for withdrawing tokens.

6 Benchmarks

In this section we give some benchmarks of our currently implemented prototype⁷. In all the benchmarks we use the 3-party replicated secret sharing protocol [MR18]. As commitments we use hash-based commitments using the Poseidon2 hash function [GKS23]. Concretely, we use Poseidon2 with an internal statesize of $t = 2$ in feed-forward mode of operation. In other words, a commitment to x with randomness r is computed as $C = \text{Poseidon2}.\text{permute}(x + d, r)[0] + x$, where d is a domain separator. We use the Noir⁸ language to write the ZK circuits, use a compiler⁹ to translate it to an R1CS system and use Co-Circom¹⁰ to create a Groth16 [Gro16] proof in MPC. Thereby, Groth16 is instantiated over the BN254 pairing friendly curve.

In all benchmarks in this section, we give numbers for different subprotocols. Thereby, **No Proofs** refers to all MPC computations without verifiability. In other words, only the data structure is updated and commitments are computed in MPC. On the contrary, **Full** refers to the full MPC protocol including the CoGroth16 proof. Furthermore, we give benchmarks for the two major parts involved in **Full**, namely **WitExt** (witness extension) which includes modifying the data structure and creating the witness for R1CS constraint system in MPC¹¹, and **CoGroth16** which translates the R1CS constraint system (including secret shared witness) into a zero knowledge proof.

6.1 Single Benchmarks

In this section we give benchmarks for processing one individual deposit, withdraw, and transfer actions. Runtimes are given, in Table 1, the corresponding throughput in actions per second in Table 2, while Table 3 gives the communication between the parties. All benchmarks in this section were performed on a co-located setup where each party is a m7a.4xlarge AWS instance with 16 CPU cores and a network bandwidth of 12.5 GBit/s.

Table 1. Runtime in ms for all the steps involved in deposit, withdraw and transfer.

Protocol	No Proofs	WitExt	CoGroth16	Full
Deposit	4.4	5.9	8.1	14.0
Withdraw	4.1	6.2	8.6	14.8
Transfer	4.3	6.9	12.3	19.2

⁷https://github.com/TaceoLabs/private_deposit

⁸<https://noir-lang.org/>

⁹<https://github.com/TaceoLabs/CoNoir-to-R1CS>

¹⁰<https://github.com/TaceoLabs/co-snarks/tree/main/co-circom>

¹¹This also implicitly creates the commitments.

Table 2. Throughput in 1/s for all the steps involved in deposit, withdraw and transfer.

Protocol	No Proofs	WitExt	CoGroth16	Full
Deposit	227.3	169.5	123.5	71.4
Withdraw	243.9	161.3	116.3	67.6
Transfer	232.6	144.9	81.3	52.1

Table 3. Communication in kB for all the steps involved in deposit, withdraw and transfer.

Protocol	No Proofs	WitExt	CoGroth16	Full
Deposit	35.1	37.8	0.2	38.0
Withdraw	35.1	41.0	0.2	41.2
Transfer	58.1	64.3	0.2	64.5

The tables show, that processing an individual action is reasonably fast, where a singlethreaded execution without proof takes roughly 4 ms, which gets increased to 19 ms for the transfer action executed with ZK proof.¹² Furthermore, communication is less than 65 kB in any case. This leads to a throughput of 230 wihtout proof and about 50 with proof.

6.2 Batched Benchmarks

In this section we give benchmarks for a multithread and batched version of the protocol (see Section 4.4). Concretely, we batch 96 transaction to match the cores on one of the used benchmark machines. Thereby, **No Proofs** and **WitExt** compute each individual transaction in a separate thread, whereas **CoGroth16** produces on ZK proof which verifies 96 transactions internally.¹³ The benchmarks are given in Table 4 for two different machine setups. Furthermore, Table 5 gives the corresponding throughput of transactions per second.

For the two machine setup we benchmark the case where each MPC party is one m7a.4xlarge AWS instance with 16 CPU cores and a network bandwidth of 12.5 GBit/s (same as for the benchmarks in the previous section), and the case where each MPC party is a m7a.24xlarge instance with a bandwidth of 37.5 GBit/s. The latter machine has 96 CPUs which matches the batch size we benchmark.

¹²Witness extension is singlethreaded, while CoGroth16 uses multithreading internally.

¹³CoGroth16 uses multithreading internally to compute the multiscalar multiplications more efficiently.

Table 4. Runtime in `ms` for all the steps involved in a batch of 96 transfers.

Machine	No Proofs	WitExt	CoGroth16	Full
m7a.4xlarge	22.9	197.8	595.8	793.6
m7a.24xlarge	15.3	185.9	273.0	458.9

Table 5. Throughput in `1/s` for all the steps involved in a batch of 96 transfers.

Machine	No Proofs	WitExt	CoGroth16	Full
m7a.4xlarge	4192.1	485.3	161.1	121.0
m7a.24xlarge	6274.5	516.4	351.6	209.2

Notice that multithreading has a profound effect on the performance. A batch of 96 transaction requires between 500 `ms` and 800 `ms`, depending on machine setup, for modifying the data structure including ZK proof. Without proof the runtime is roughly 20 `ms`. Throughput, thus is increased to roughly 200 transactions per second with proof, and 6000 for a non-verifiable variant.

Notice, that there is not much difference between the runtime of the witness extension (`WitExt`) in the two different machine setups. We believe that the small difference is mostly due to the larger network bandwidth of the m7a.24xlarge instances, and that the performance of the witness extension is mostly network bound.

Proof generation, on the other hand, is clearly CPU bound, where the increase of number of CPUs from 16 to 96 improves performance by more than two-fold.

6.3 Solidity Gas Costs

To demonstrate the feasibility of our approach, we implement a small solidity smart contract with the capabilities described in this paper. When deployed, a client can register deposits, withdraws, and transfers on chain which gets stored in an action queue. An MPC network queries this queue and processes a batch of transactions at a time. Once done, the MPC network posts the ZK proof and the transactions are counted as fulfilled. In this implementation, client-to-MPC-network communication (i.e., sending over the secret shared amount and randomness for a transfer) is done via the smart contract by adding ciphertexts of the shares to the action queue when registering a transfer.

In Table 6 we give the gas costs of each the 4 main solidity functions (`deposit`, `withdraw`, `transfer` for the client and `processMPC` for the MPC network) when executed on the Base

Sepolia testnet.¹⁴ In this table, `processMPC` is implemented to process a batch of 50 transactions per proof. We further give the gas cost in GWEI assuming a gas price of 0.017 GWEI.¹⁵

Table 6. Amount of gas used when deployed on Base Sepolia Testnet and gas price assuming a cost of 0.017 GWEI.

Function	Gas	Gas cost [GWEI]
<code>Deposit</code>	150869	2564.773
<code>Withdraw</code>	136838	2326.246
<code>Transfer</code>	336856	5726.552
<code>ProcessMPC (50)</code>	3311538	56296.146

7 Issues and Future Work

In this section we discuss some issues of the currently described protocol and potential future work.

7.1 Implications of Batching Transactions

A prototype implemented from the batched version of our protocol (see Section 4.4 and Section 6.2) has the advantage that a batch of transaction (e.g., 96) only creates 1 proof, reducing on-chain gas fees for proof verification. However, it comes with some drawbacks as well.

First, when creating a ZK proof for a single transaction, the smart contract and the MPC network immediately know if the transaction was valid or faulty by verifying the proof. In a batched version of the proof, however, if one transaction fails, the whole batch fails. Currently, there is also no method implemented to detect which of the transactions in a batch were valid or faulty. We think of two potential solutions to this issue.

The first solution is checking in MPC whether the range checks in the zero knowledge proofs (see Listing 3) would be valid or not and open the resulting valid-bit. Based on this bit, the MPC network would then decide to not include the transaction in a batch. To cleanup the on-chain action queue, the MPC network would then call a removal function exposed by the smart contract to remove the faulty transaction from the queue. While this solution is simple and should induce only small performance overhead, it comes with the disadvantage that the existence of a removal function for the on-chain action queue gives the MPC network the possibility to ignore transactions and censor specific user addresses.

¹⁴A sample deployment where each function is executed once can be found here: <https://sepolia.basescan.org/address/0xf55dfe8242ec117f44acd932adb5c233b582a04a>.

¹⁵Taken from <https://basescan.org/gastracker> for Base mainnet at the time of writing.

The second solution is modifying the zero knowledge proofs to be valid even for faulty transactions, but outputting a bit indicating whether the transaction was valid. In other words, the range constraints are modified to output 0 if in range and 1 if not in range. This bit is an extra input to the smart contract and the proof verification, and if the bit is set, the smart contract removes the faulty action from the queue.

If the latter solution is implemented, it is probably also required to let the smart contract enforce the order of the processed transaction to prevent letting the MPC network just ignoring some transactions indefinitely.

7.2 Concurrent Transactions and Sharding

The current prototype implements multithreading by first modifying the data structure in sequence, and then spawning a thread to compute the witness extension fully in parallel. This has the advantage that we do not need to care about multiple transactions modifying the same commitments in our hashmap concurrently which would lead to faulty results. Furthermore, modifying the data structure is simple and requires no communication, so no performance is lost by doing it in sequence.

This also has the advantage that we can scale up our solution with sharding computations to multiple servers. First one thread takes care of modifying the data structure and then delegates the witness extension and proof generation to a different MPC setup.

The downside of such an approach, however, is that if the MPC network detects faulty transactions, it needs to rewind all the transactions that came after the faulty one, since they might have been using a faulty commitment as input. This would not be the case if a batch only contains a disjoint set of senders and receivers.

7.3 Hiding Involved Parties

As discussed in Section 2, we do not hide sender and receiver addresses during private transactions to keep compliance simple. Thus, future work can investigate a fully private transaction solution which hides sender and receiver, alongside balances and transfer amounts, as well. This raises two issues: First, on the technical side, accessing a hashmap at private indices is more involved in MPC and might require the usage of oblivious RAM techniques and similar. This might also require to only keep a sparse commitment (e.g., from a Merkle-tree) on chain, making zero-knowledge proofs more involved as well. Second, a thorough investigation of how to keep compliance might be required. This might include a know-your-customer step during a registration and proofs that no blacklisted funds have been received.

References

- [Yao82] A. C.-C. Yao, “Protocols for Secure Computations (Extended Abstract),” in *FOCS*, IEEE Computer Society, 1982, pp. 160–164.

- [Gen09] C. Gentry, “Fully homomorphic encryption using ideal lattices,” in *STOC*, ACM, 2009, pp. 169–178.
- [OB22] A. Ozdemir and D. Boneh, “Experimenting with Collaborative zk-SNARKs: Zero-Knowledge Proofs for Distributed Secrets,” in *USENIX Security Symposium*, USENIX Association, 2022, pp. 4291–4308.
- [Bow+20] S. Bowe, A. Chiesa, M. Green, I. Miers, P. Mishra, and H. Wu, “ZEXE: Enabling Decentralized Private Computation,” in *SP*, IEEE, 2020, pp. 947–964.
- [Zys+25] H.-T. U. C. T. F. D. G. Zyskind, A. Chiesa, D. Zarchy, M. Leibovich, and C. Peikert, “High-Throughput Universally Composable Threshold FHE Decryption,” *IACR Cryptol. ePrint Arch.*, p. 1781, 2025.
- [Bea91] D. Beaver, “Efficient Multiparty Protocols Using Circuit Randomization,” in *CRYPTO*, in Lecture Notes in Computer Science, vol. 576. Springer, 1991, pp. 420–432.
- [Chi+20] I. Chillotti, N. Gama, M. Georgieva, and M. Izabachène, “TFHE: Fast Fully Homomorphic Encryption Over the Torus,” *J. Cryptol.*, vol. 33, no. 1, pp. 34–91, 2020.
- [MR18] P. Mohassel and P. Rindal, “ABY3: A Mixed Protocol Framework for Machine Learning,” in *CCS*, ACM, 2018, pp. 35–52.
- [GKS23] L. Grassi, D. Khovratovich, and M. Schafneger, “Poseidon2: A Faster Version of the Poseidon Hash Function,” in *AFRICACRYPT*, in Lecture Notes in Computer Science, vol. 14064. Springer, 2023, pp. 177–203.
- [Gro16] J. Groth, “On the Size of Pairing-Based Non-interactive Arguments,” in *EUROCRYPT (2)*, in Lecture Notes in Computer Science, vol. 9666. Springer, 2016, pp. 305–326.